

AY: 2024-2025  
MIDTERM | AI-ECUE221  
Apr. 2025

M1-S2: Dept. of Electrical Engineering  
Teacher: A. Mhamdi  
Time Limit: 1h

This document contains 7 pages numbered from 1/7 to 7/7. As soon as it is handed over to you, make sure it is complete. The 3 tasks are independent and can be treated in the order that suits you.

The following rules apply:

- ❶ A handwritten double-sided A4 sheet is permitted.
- ❷ Any electronic material, except basic calculator, is prohibited.
- ❸ Mysterious or unsupported answers will not receive full credit.
- ❹ Round results to the nearest thousandth (i.e., third digit after the decimal point).
- ❺ Task N°3: Each correct answer will grant a mark with no negative scoring.

### Task N°1

⌚ 30mn | (8 points)

We consider the following dataset ( $x_1 = 1$ ):

Table 1: Housing energy efficiency prediction.

House Area ( $x_2$ )	Insulation Thickness ( $x_3$ )	Window ( $x_4$ )	Quality	Energy Consumption ( $y$ )
100	5	3		200
120	6	4		180
150	4	2		220
200	7	5		150
180	5	4		170

The output  $y$  can be expressed as a linear function of house characteristics (i.e.,  $y = \mathbf{x}^T \cdot \boldsymbol{\theta}$ ).  
An estimate of the parameter vector  $\boldsymbol{\theta}$  is given by:

$$\hat{\boldsymbol{\theta}} \approx [266.667, -0.076, 3.806, -25.520]^T$$

- (a) (2 points) Predict energy consumption for a house with: Area=170, Window Quality=4, and Insulation Thickness=6.

$$\begin{aligned}\hat{y} &= \underbrace{[1, 170, 6, 4]}_{x^T} \cdot \underbrace{[266.667, -0.076, 3.806, -25.520]^T}_{\hat{\theta}} \\ &\approx 174.503\end{aligned}$$

(b) (6 points) Compute each of the following metrics:

MAE

RMSE

MAPE

R-squared

Given the value of  $\hat{\theta}$ , we can compute the predicted output:

$$\hat{y} = x \cdot \hat{\theta} \approx [201.537, 178.303, 219.451, 150.509, 169.937]^T$$

The error vector  $\varepsilon$  is:

$$\varepsilon = y - \hat{y} \approx [-1.537, 1.697, 0.549, -0.509, 0.063]^T$$

**Mean Absolute Error (MAE)**

$$\begin{aligned}\text{MAE} &= \frac{1}{5} \sum_{i=1}^5 |\varepsilon_i| \\ &\approx 0.871\end{aligned}$$

**Root Mean Squared Error (RMSE)**

$$\begin{aligned}\text{RMSE} &= \sqrt{\frac{1}{5} \sum_{i=1}^5 \varepsilon_i^2} \\ &\approx 1.078\end{aligned}$$

**Mean Absolute Percentage Error (MAPE)**

$$\begin{aligned}\text{MAPE} &= \frac{1}{5} \sum_{i=1}^5 \left| \frac{\varepsilon_i}{y_i} \right| \cdot 100\% \\ &\approx 0.467\%\end{aligned}$$

### R-squared

$$R^2 = 1 - \frac{\sum_{i=1}^5 (y_i - \hat{y}_i)^2}{\sum_{i=1}^5 (y_i - \bar{y})^2}$$
$$\approx 0.998$$
$$\bar{y} = \frac{200 + 180 + 220 + 150 + 170}{5} = 184.0$$

ANSWERS

AY: 2024-2025

M1-S2: Dept. of Electrical Engineering

MIDTERM | AI-ECUE221

Apr. 2025

Teacher: A. Mhamdi

Full Name: .....

ID: .....

Class: RAIA1 .....

Room: .....

Time Limit: 1h

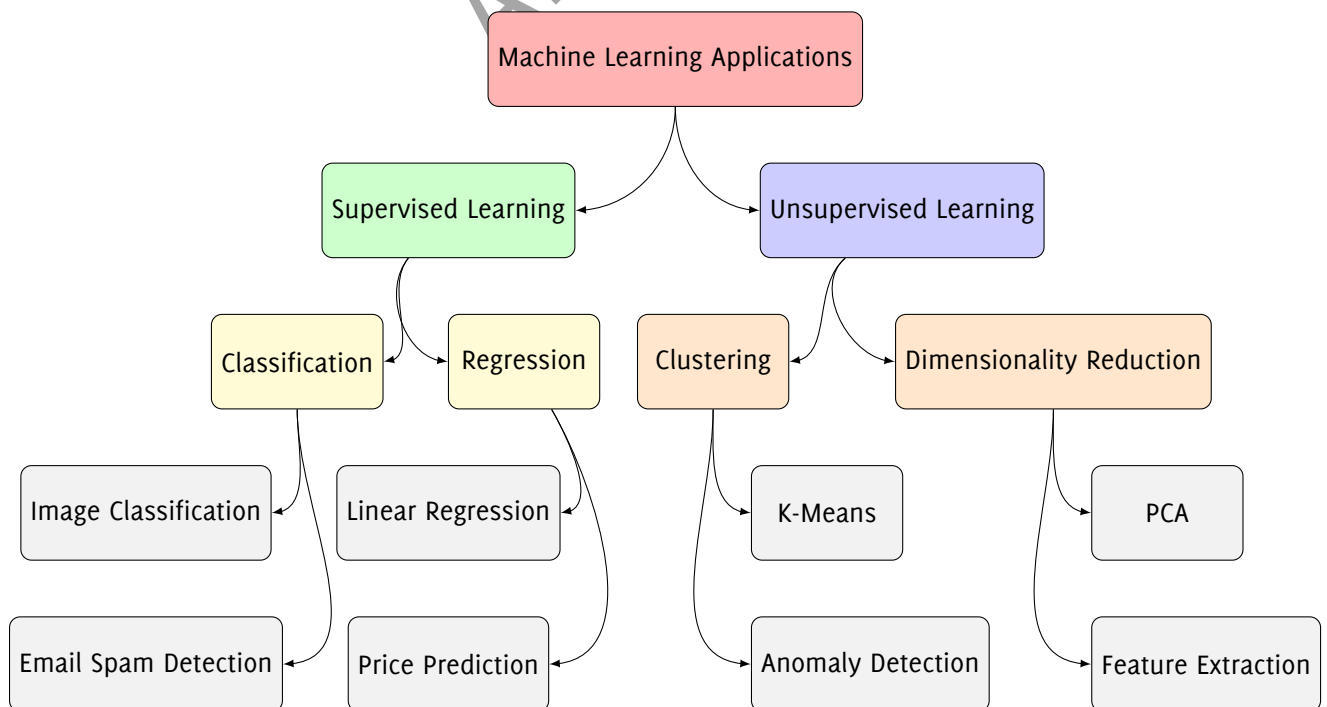
ANSWER SHEET

Task N°2

⌚ 15mn | (3½ points)

Categorize the following machine learning applications into their appropriate learning paradigms by placing each one in the correct branch:

Anomaly Detection	Feature Extraction	PCA
Classification	Image Classification	Email Spam Detection
Unsupervised Learning	Linear Regression	Price Prediction
Supervised Learning	Clustering	Dimensionality Reduction
K-Means	Regression	



DO NOT WRITE ANYTHING HERE

✂

**Task N°3**

⌚ 15mn | (8½ points)

- (a) (½ point) Which technique helps reduce overfitting in linear regression?
- ☐ Increasing model complexity
  - ☐ Adding polynomial features
  - ✓ ☒ Using regularization (L1/L2)
  - ☐ Decreasing training data
- (b) (½ point) The cost function for linear regression is typically:
- ☐ Mean Absolute Error (MAE)
  - ✓ ☒ Mean Squared Error (MSE)
  - ☐ Cross-Entropy Loss
  - ☐ Hinge Loss
- (c) (½ point) Gradient descent is primarily used for:
- ☐ Data visualization
  - ☐ Feature engineering
  - ☐ Data cleaning
  - ✓ ☒ Optimizing model parameters
- (d) (½ point) What is the purpose of splitting data into training and test sets?
- ☐ To increase model complexity
  - ✓ ☒ To evaluate generalization performance
  - ☐ To reduce feature dimensions
  - ☐ To compute gradients
- (e) (½ point) What is MLJ's primary approach to handling different data types in machine learning workflows?
- ✓ ☒ Uses Julia's native Tables.jl interface for data compatibility
  - ☐ Only works with CSV files
  - ☐ Requires manual type conversion
  - ☐ Requires conversion to NumPy arrays
- (f) (½ point) What is the recommended way to handle missing values in MLJ?

DO NOT WRITE ANYTHING HERE

✂

- ☐ Only works with complete datasets
- ☐ Automatically drops rows with missing values
- ☐ Requires manual imputation before model fitting
- ✓ Provides built-in transformers like FillImputer

(g) ( $\frac{1}{2}$  point) What is MLJ's recommended method for feature scaling?

- ☐ No scaling support
- ☐ Manual z-score calculation
- ✓ Standardizer transformer from MLJModels
- ☐ External Python preprocessing

(h) (1 point) Which feature makes MLJ particularly suitable for reproducible data preprocessing?

- ☐ Automatic hyperparameter optimization
- ✓ Composable, persistent pipeline blueprints
- ☐ Exclusive use of deep learning models
- ☐ Integration with JavaScript visualizations

(i) (1 point) What does this code do?

```
1 using DataFrames
2 df = DataFrame(A=1:5, B=11:15)
3 df = transform(df, :A => (x -> x.^2) => :A_squared)
```

- ☐ Filters rows where  $A > 3$
- ☐ Replaces column A with its squared values
- ✓ Creates a new column A\_squared containing squares of column A
- ☐ Groups data by column A

(j) (1 point) What is the purpose of this code?

```
1 using DataFrames, MLJ
2 df = DataFrame(X=[1, missing, 3], Y=[5, 6, missing])
3 df = coalesce.(df, 0)
```

- ☐ Drops all rows with missing values

DO NOT WRITE ANYTHING HERE

✂

✓ Replaces missing values with 0 in all columns

- ☐ Imputes missing values using column means
- ☐ Creates dummy variables for missing entries

(k) (1 point) What does this pipeline accomplish?

```
1 using MLJ
2 LR = @load LinearRegressor pkg=GLM
3 pipe = FillImputer() |> Standardizer() |> LR()
```

- ☐ One-hot encodes features, then fits a linear model
- ☐ Imputes missing values with 0, skips scaling, then fits a GLM
- ✓ Fills missing values with column means, standardizes features, then fits a linear regression
- ☐ Removes rows with missing values, normalizes features, then fits a decision tree

(l) (1 point) Which code snippet will run without errors?

```
1 # Option 1
2 mutable struct Sensor
3     id::Int
4     readings::Vector{Float64}
5 end
6 s = Sensor(1, [1.0, 2.0])
7 push!(s.readings, 3.0)
8
9 # Option 2
10 struct Sensor
11     id::Int
12     readings::Vector{Float64}
13 end
14 s = Sensor(1, [1.0, 2.0])
15 push!(s.readings, 3.0)
```

- ☐ Neither works
- ☐ Only Option 1 works
- ☐ Only Option 2 works
- ✓ Both work